

Pune Vidyarthi Griha College of Engineering
and Technology Pune

DIGITAL SIGNAL PROCESSING

Application Assignment

Speech De-noising Using Deep Learning

Hybrid Group No: 5

Group Member 1: Mohammad Shahbaz Alam

Group Member 2: Devidas Ingale

Group Member 3: Avdhoot Patil

Guided By: Dr Vikram Gadre Sir

June 20, 2021

1. Title
2. Software/Platform used
3. Steps to install and run the code
4. Stepwise windows screenshots
5. Block Diagram
6. Algorithm used
7. Observation
8. Conclusion
9. Code Links

1) Title

Speech Denoising Using Deep learning

2) Software/Platform used

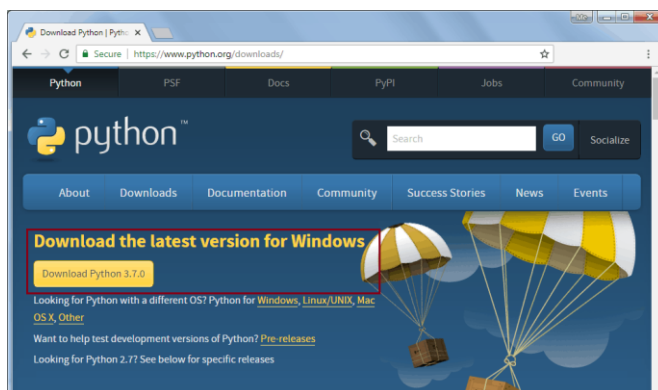
Coding: Python in Jupyter Notebook IDE

Neural Network made by: TensorFlow 2

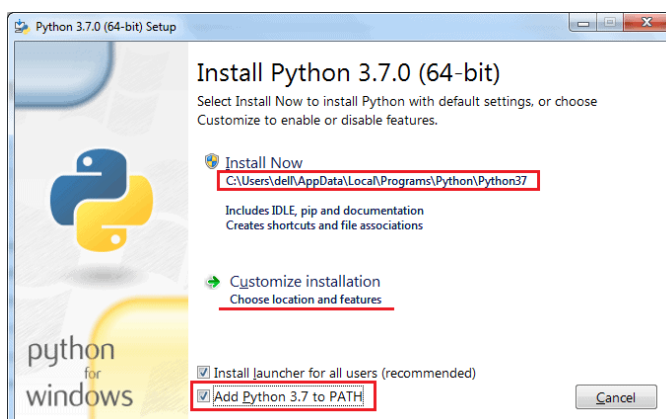
Libraries: NumPy, Librosa, Soundfile etc.

3) Steps to install & run the code

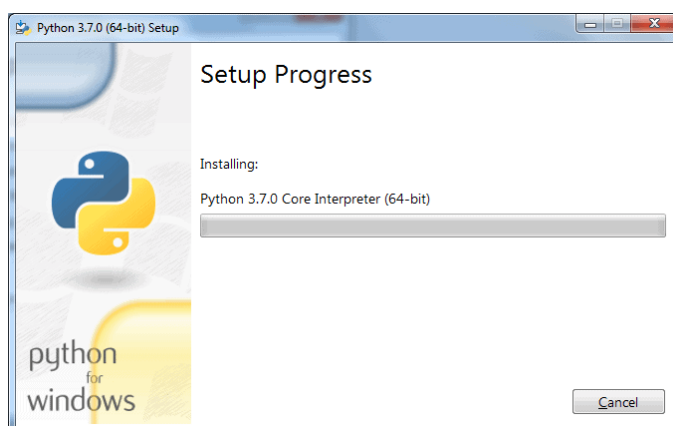
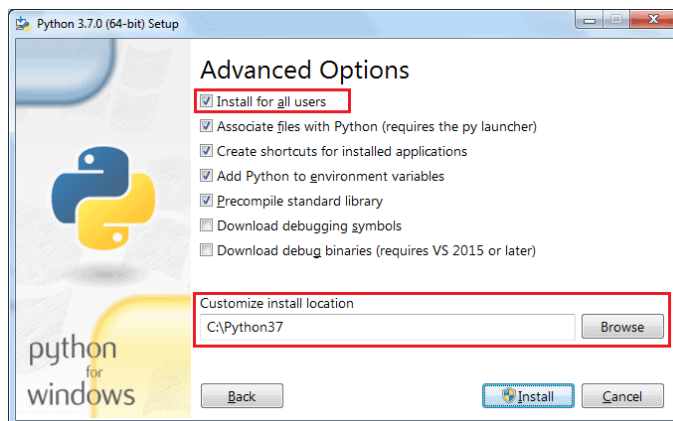
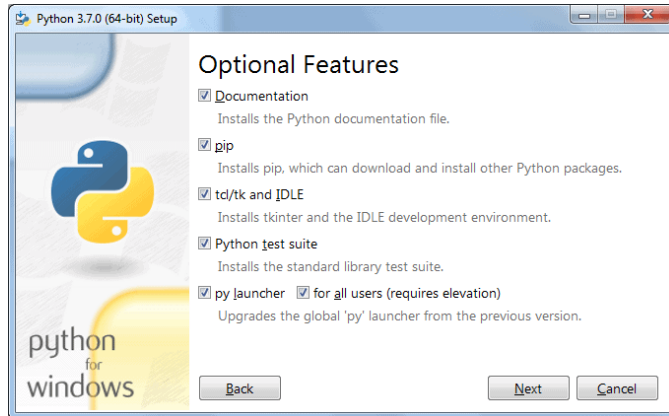
- 1) Download python in given link: <https://www.python.org/downloads/>



- Install python



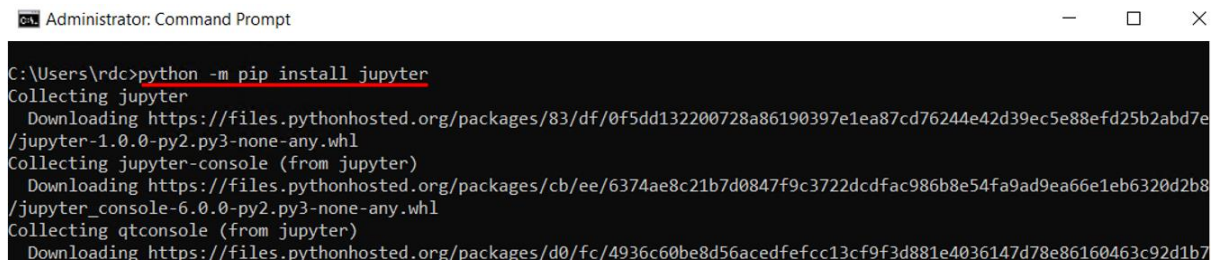
- Click all optional features including pip



- Python is installed

2) Now installing Jupyter notebook using pip in command prompt

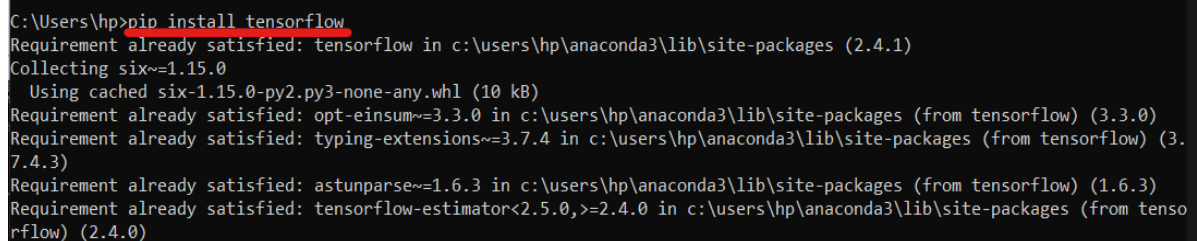
```
python -m pip install jupyter
```



```
Administrator: Command Prompt
C:\Users\rndc>python -m pip install jupyter
Collecting jupyter
  Downloading https://files.pythonhosted.org/packages/83/df/0f5dd132200728a86190397e1ea87cd76244e42d39ec5e88efd25b2abd7e/jupyter-1.0.0-py2.py3-none-any.whl
Collecting jupyter-console (from jupyter)
  Downloading https://files.pythonhosted.org/packages/cb/ee/6374ae8c21b7d0847f9c3722dcdfac986b8e54fa9ad9ea66e1eb6320d2b8/jupyter_console-6.0.0-py2.py3-none-any.whl
Collecting qtconsole (from jupyter)
  Downloading https://files.pythonhosted.org/packages/d0/fc/4936c60be8d56acedfefcc13cf9f3d881e4036147d78e86160463c92d1b7
```

3) Install TensorFlow 2 using pip in command prompt:

```
pip install tensorflow
```



```
C:\Users\hp>pip install tensorflow
Requirement already satisfied: tensorflow in c:\users\hp\anaconda3\lib\site-packages (2.4.1)
Collecting six~1.15.0
  Using cached six-1.15.0-py2.py3-none-any.whl (10 kB)
Requirement already satisfied: opt-einsum~=3.3.0 in c:\users\hp\anaconda3\lib\site-packages (from tensorflow) (3.3.0)
Requirement already satisfied: typing-extensions~=3.7.4 in c:\users\hp\anaconda3\lib\site-packages (from tensorflow) (3.7.4.3)
Requirement already satisfied: astunparse~=1.6.3 in c:\users\hp\anaconda3\lib\site-packages (from tensorflow) (1.6.3)
Requirement already satisfied: tensorflow-estimator<2.5.0,>=2.4.0 in c:\users\hp\anaconda3\lib\site-packages (from tensorflow) (2.4.0)
```

4) Installing other libraries like librosa, NumPy, soundfile etc using pip in command prompt:

```
pip install librosa
```

```
pip install numpy
```

```
pip install soundfile
```

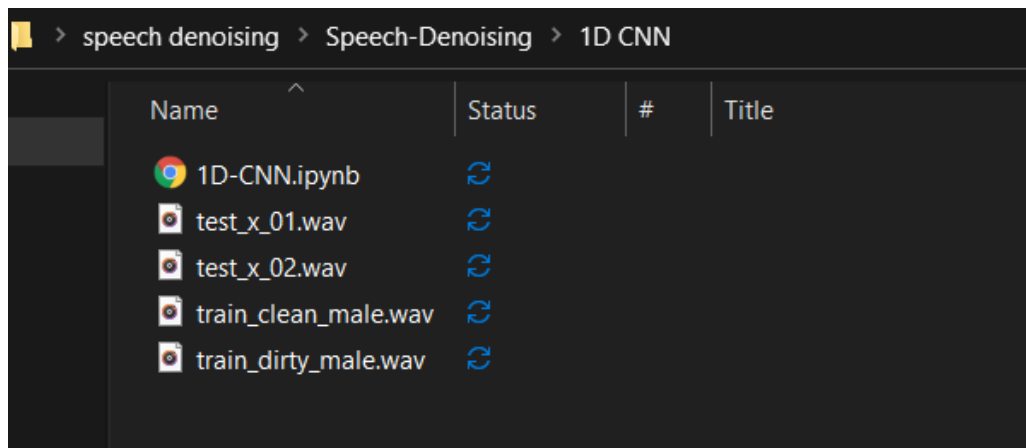
```
pip install matplotlib
```

```
pip install scipy
```

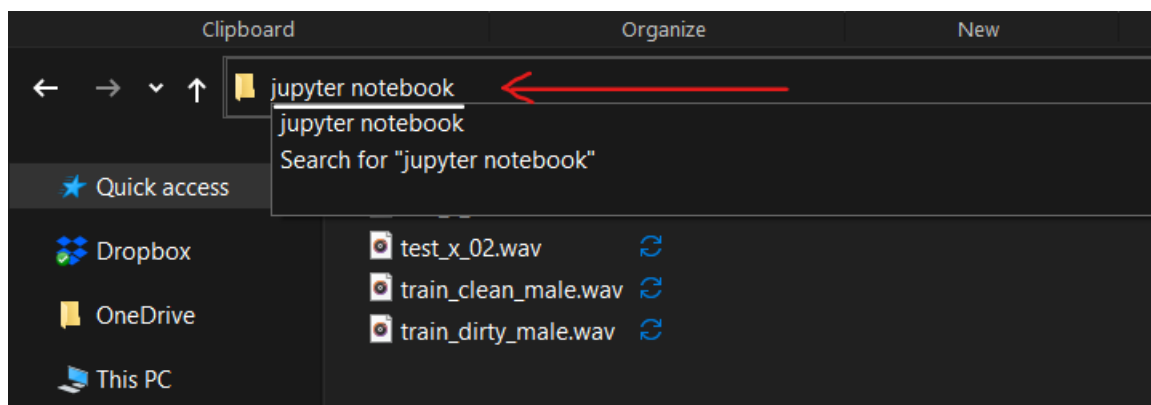
```
pip install wavefile
```

```
pip install ipython
```

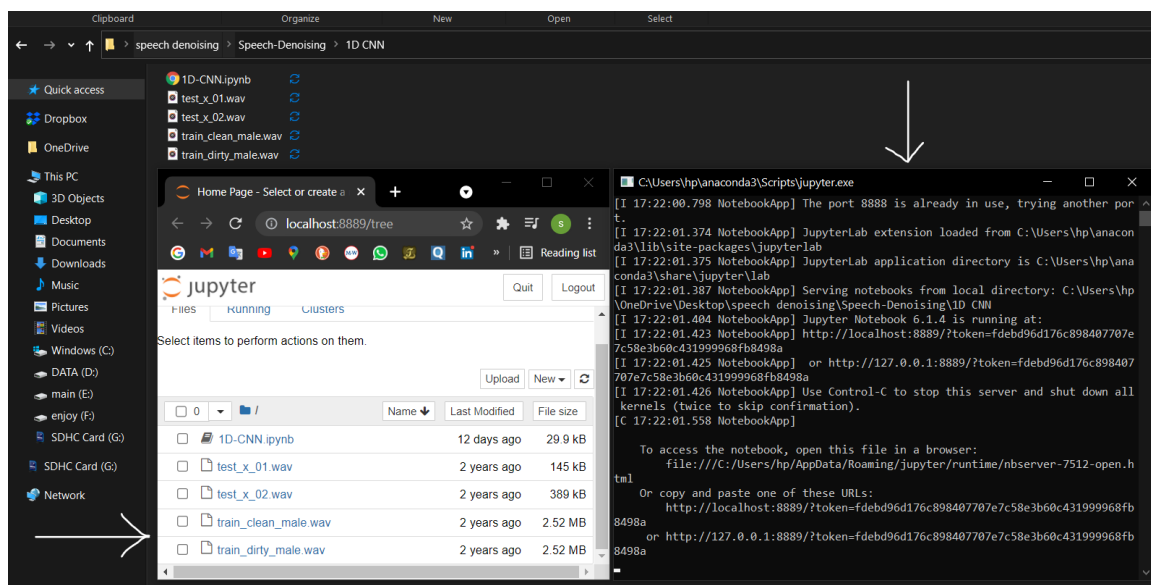
Open File Containing 1D CNN jupyter with test & training files



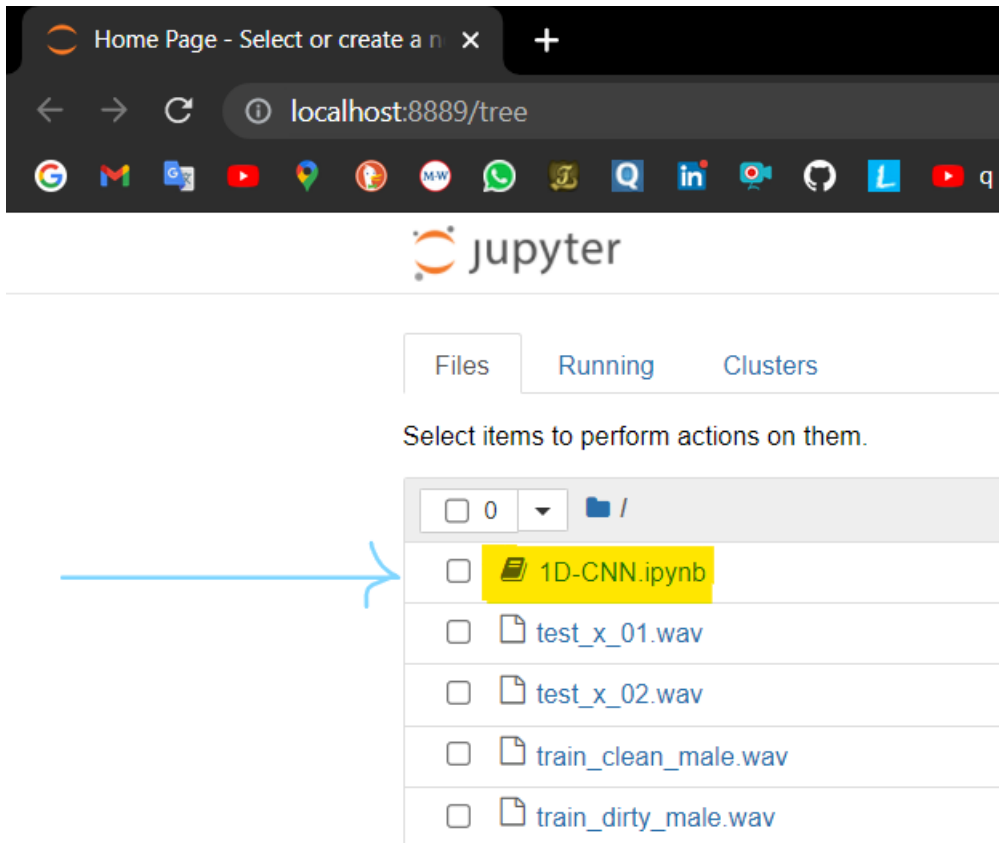
In location bar, write jupyter notebook, then press “enter”. It opens terminal



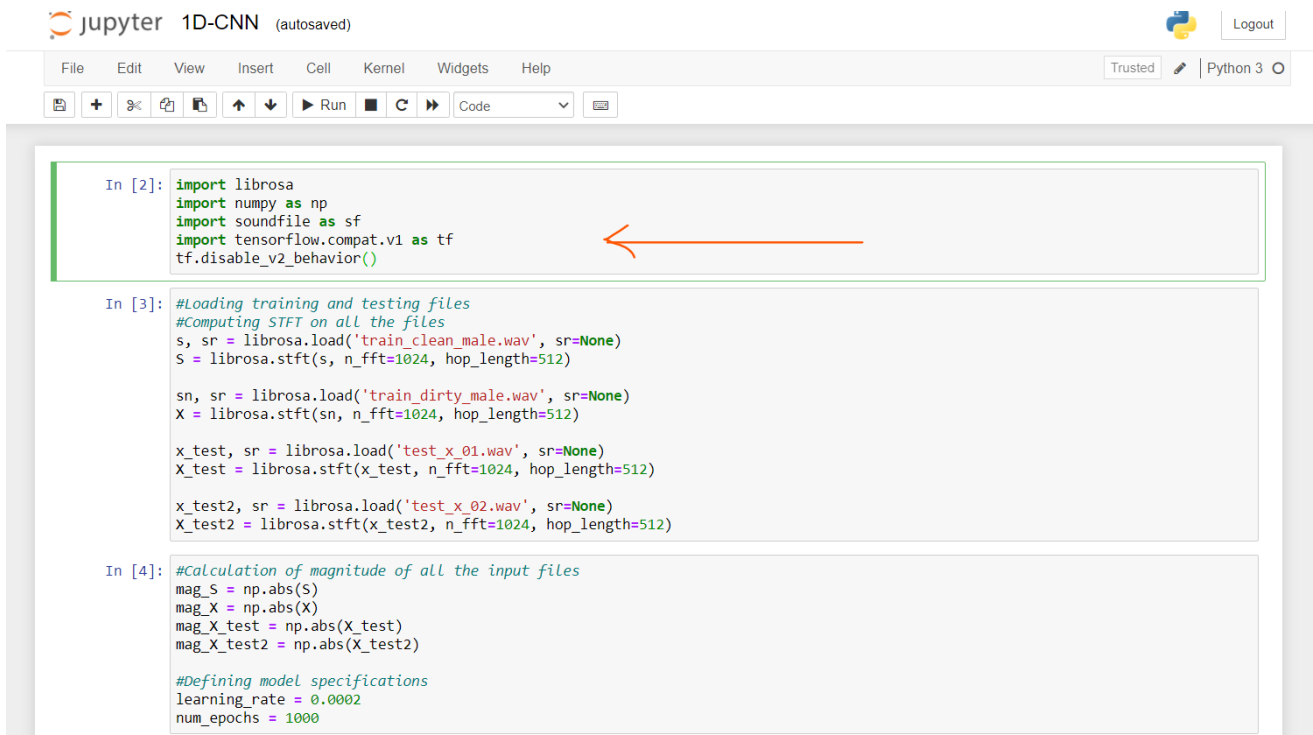
This will open terminal and jupyter notebook



Open 1D-CNN.pynb file



Click in this block and press “shift + Enter” to run this block and continue till training of file



Run by pressing “shift + Enter” till this block starts training of audio files

Jupyter 1D-CNN Last Checkpoint: 28 minutes ago (autosaved)

```
In [6]: output = getModel(input)
#Defining the loss function along with its optimizer
loss = tf.reduce_mean(tf.square(output - labels))
train_step = tf.train.AdamOptimizer(learning_rate).minimize(loss)

sess = tf.InteractiveSession()
sess.run(tf.global_variables_initializer())

count = 0
batch_size = 100
flag = True

while flag:
    size = 0
    #Mini batching with the given batch size
    for i in range(0, 2459, batch_size):
        size += batch_size
        if size <= 2459:
            batch_x = mag_X[:,i : size]
            batch_y = mag_S[:,i : size]
        else:
            batch_x = mag_X[:,i : 2459]
            batch_y = mag_S[:,i : 2459]

        feed_dict = {input: batch_x.T, labels: batch_y.T}
        train_step.run(feed_dict=feed_dict)

    if count%2 == 0:
        loss_calc = loss.eval(feed_dict=feed_dict)
        print("Epoch %d, loss %g"%(count, loss_calc))
```

Now, training starts

```
#Once all the epochs are completed, training is stopped
if count >= num_epochs:
    flag = False

count+=1
```

C:\Users\hp\anaconda3\lib\site-packages\tensorflow\python\keras\legacy_tf_layers\conv1d is deprecated and will be removed in a future version. Please use `tf.keras.layers.Conv1D` instead. warnings.warn('tf.layers.conv1d is deprecated and')

C:\Users\hp\anaconda3\lib\site-packages\tensorflow\python\keras\engine\base_layer_v1.py:100: DeprecationWarning: `layer.apply` is deprecated and will be removed in a future version. Please use `layer.__call__` method instead. warnings.warn('layer.apply is deprecated and')

C:\Users\hp\anaconda3\lib\site-packages\tensorflow\python\keras\legacy_tf_layers\pooling1d.py:100: DeprecationWarning: `tf.layers.max_pooling1d` is deprecated and will be removed in a future version. Please use `tf.keras.layers.MaxPooling1D` instead. warnings.warn('tf.layers.max_pooling1d is deprecated and')

C:\Users\hp\anaconda3\lib\site-packages\tensorflow\python\keras\legacy_tf_layers\core.py:100: DeprecationWarning: `tf.layers.flatten` is deprecated and will be removed in a future version. Please use `tf.keras.layers.Flatten` instead. warnings.warn('tf.layers.flatten is deprecated and')

C:\Users\hp\anaconda3\lib\site-packages\tensorflow\python\keras\legacy_tf_layers\core.py:100: DeprecationWarning: `tf.layers.dense` is deprecated and will be removed in a future version. Please use `tf.keras.layers.Dense` instead. warnings.warn('tf.layers.dense is deprecated and')

```
Epoch 0, loss 0.0127375
Epoch 2, loss 0.00856482
Epoch 4, loss 0.00694965
Epoch 6, loss 0.00555794
Epoch 8, loss 0.00429112
Epoch 10, loss 0.00355224
Epoch 12, loss 0.00311731
Epoch 14, loss 0.00281695
Epoch 16, loss 0.0025776
Epoch 18, loss 0.00236361
Epoch 20, loss 0.00221721
Epoch 22, loss 0.00210539
Epoch 24, loss 0.00199769
```


Here, training completes. So again, start running next blocks

```
Epoch 958, loss 0.000398425
Epoch 960, loss 0.000393677
Epoch 962, loss 0.000400141
Epoch 964, loss 0.000417317
Epoch 966, loss 0.000420429
Epoch 968, loss 0.000438755
Epoch 970, loss 0.000455793
Epoch 972, loss 0.000457004
Epoch 974, loss 0.000448257
Epoch 976, loss 0.000437344
Epoch 978, loss 0.000409197
Epoch 980, loss 0.000392719
Epoch 982, loss 0.000386491
Epoch 984, loss 0.000377681
Epoch 986, loss 0.000383855
Epoch 988, loss 0.000388827
Epoch 990, loss 0.000391897
Epoch 992, loss 0.00039956
Epoch 994, loss 0.000402182
Epoch 996, loss 0.000396761
Epoch 998, loss 0.000393115
Epoch 1000, loss 0.000389643
```

Finally, we get SNR of 17.6417 by 1D CNN

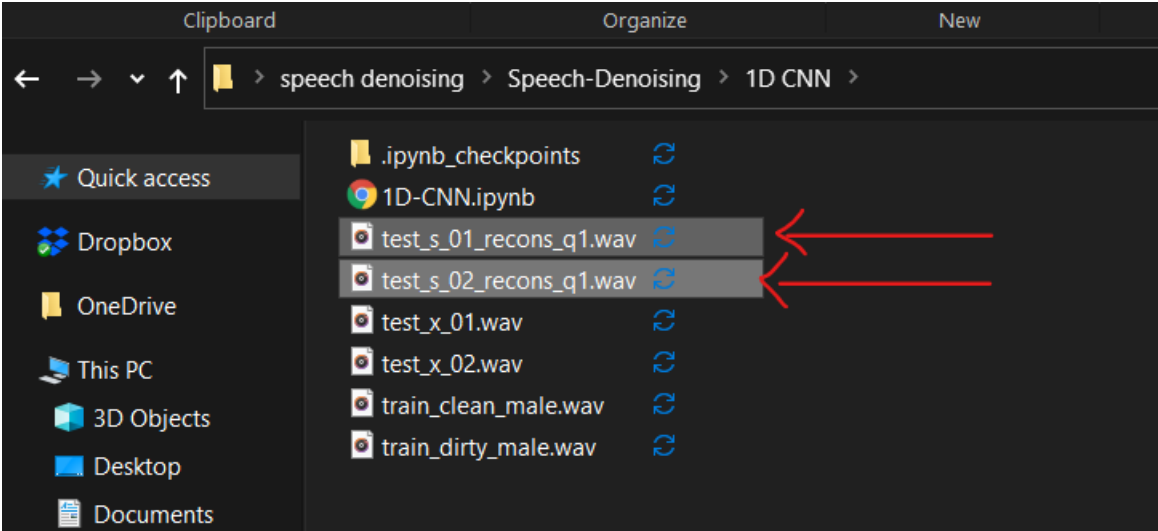
```
File Edit View Insert Cell Kernel Widgets Help

In [15]: #For testing purpose, feeding the model with train_dirty_male file
#From the output generated, reconstructing the audio file
s_hat_test3 = feedforward(mag_X.T , output)
s_hat3 = recover_sound(X, mag_X , s_hat_test3.T)
recon_sound3 = librosa.istft(s_hat3 , hop_length=512 , win_length=1024)
size_recon_sound3 = np.shape(recon_sound3)[0]

In [16]: #Once the audio file is generated, calculating the SNR value
s = s[: size_recon_sound3]
num = np.dot(s.T , s)
den = np.dot((s - recon_sound3).T,(s - recon_sound3))
SNR = 10 * np.log10(num/den)
print('Value of SNR : ' + str(SNR))

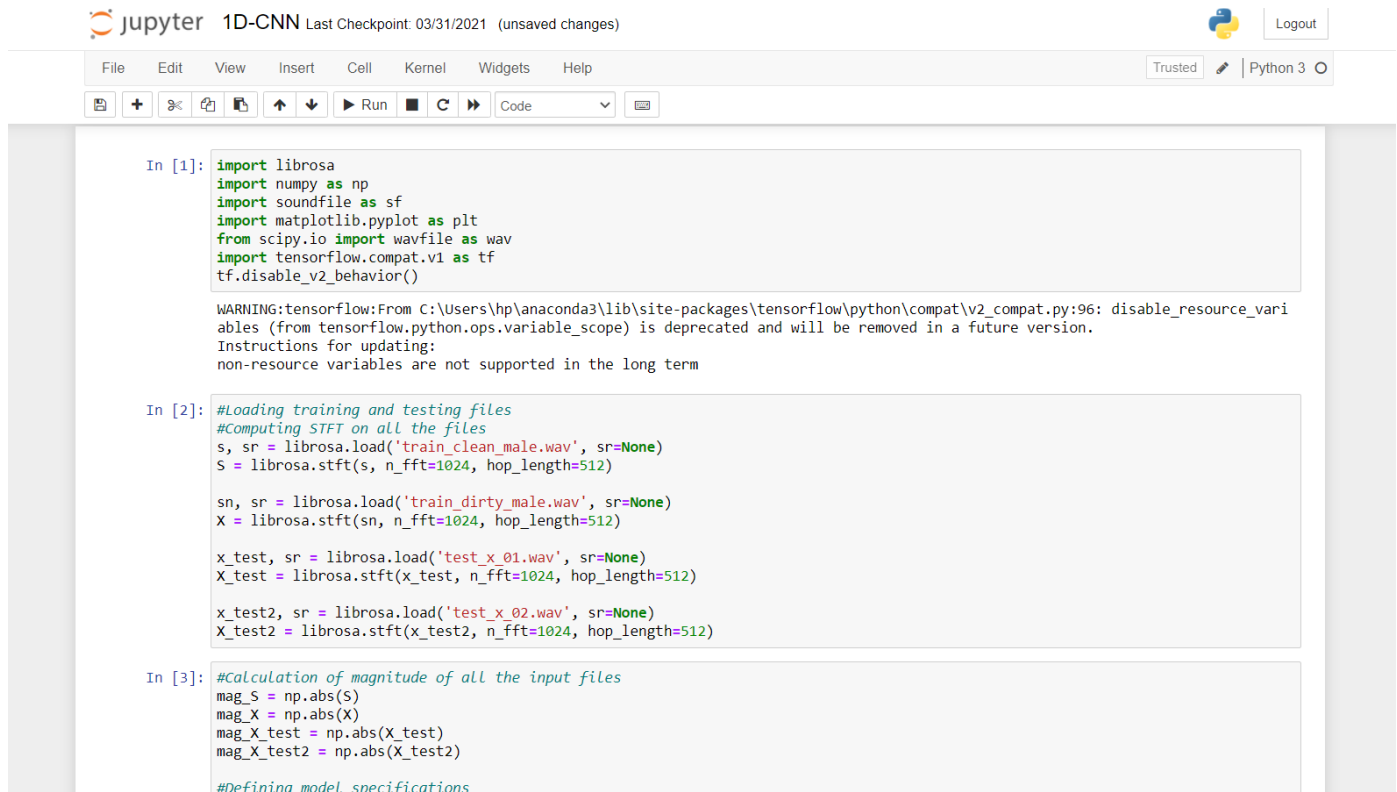
Value of SNR : 17.64175057411194
```

After running jupyter notebook we get our output 1D CNN denoised audio wav file in our given location



4) Stepwise windows screenshot

Working with jupyter notebook for speech denoising using deep learning.



The screenshot shows a Jupyter Notebook interface with the title "1D-CNN". The top bar includes the Jupyter logo, the title, and a "Logout" button. Below the title bar is a menu bar with options: File, Edit, View, Insert, Cell, Kernel, Widgets, and Help. A "Trusted" status indicator and a "Python 3" version selector are also present. The main area contains three code cells. The first cell imports various libraries: librosa, numpy, soundfile, matplotlib.pyplot, scipy.io, and tensorflow.compat.v1. The second cell loads and processes audio files, including training and testing files, and computes STFT. The third cell calculates the magnitude of the input files and defines model specifications. A warning message from tensorflow is displayed between the first and second code cells.

```
In [1]: import librosa
import numpy as np
import soundfile as sf
import matplotlib.pyplot as plt
from scipy.io import wavfile as wav
import tensorflow.compat.v1 as tf
tf.disable_v2_behavior()

WARNING:tensorflow:From C:\Users\hp\anaconda3\lib\site-packages\tensorflow\python\compat\v2_compat.py:96: disable_resource_variables (from tensorflow.python.ops.variable_scope) is deprecated and will be removed in a future version.
Instructions for updating:
non-resource variables are not supported in the long term

In [2]: #Loading training and testing files
#Computing STFT on all the files
s, sr = librosa.load('train_clean_male.wav', sr=None)
S = librosa.stft(s, n_fft=1024, hop_length=512)

sn, sr = librosa.load('train_dirty_male.wav', sr=None)
X = librosa.stft(sn, n_fft=1024, hop_length=512)

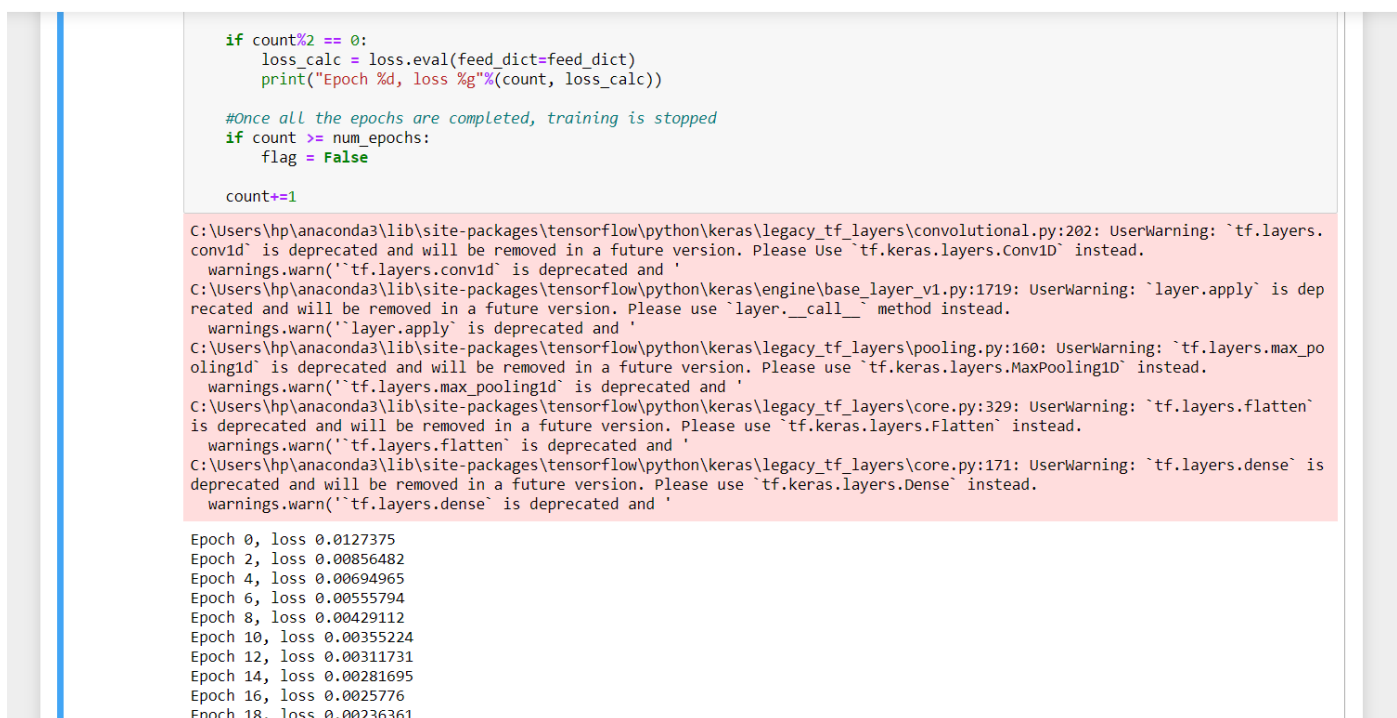
x_test, sr = librosa.load('test_x_01.wav', sr=None)
X_test = librosa.stft(x_test, n_fft=1024, hop_length=512)

x_test2, sr = librosa.load('test_x_02.wav', sr=None)
X_test2 = librosa.stft(x_test2, n_fft=1024, hop_length=512)

In [3]: #Calculation of magnitude of all the input files
mag_S = np.abs(S)
mag_X = np.abs(X)
mag_X_test = np.abs(X_test)
mag_X_test2 = np.abs(X_test2)

#Defining model specifications
```

Training files:



The screenshot shows a Jupyter Notebook interface displaying the training progress of a 1D-CNN model. The code cell shows the training loop, which prints the loss for each epoch. The output shows the loss decreasing over 18 epochs. A large block of warnings is displayed, indicating deprecated TensorFlow layers and methods. The warnings include: 'tf.layers.conv1d' is deprecated, 'tf.layers.conv1d' is deprecated, 'layer.apply' is deprecated, 'tf.layers.max_pooling1d' is deprecated, 'tf.layers.flatten' is deprecated, and 'tf.layers.dense' is deprecated.

```
if count%2 == 0:
    loss_calc = loss.eval(feed_dict=feed_dict)
    print("Epoch %d, loss %g"%(count, loss_calc))

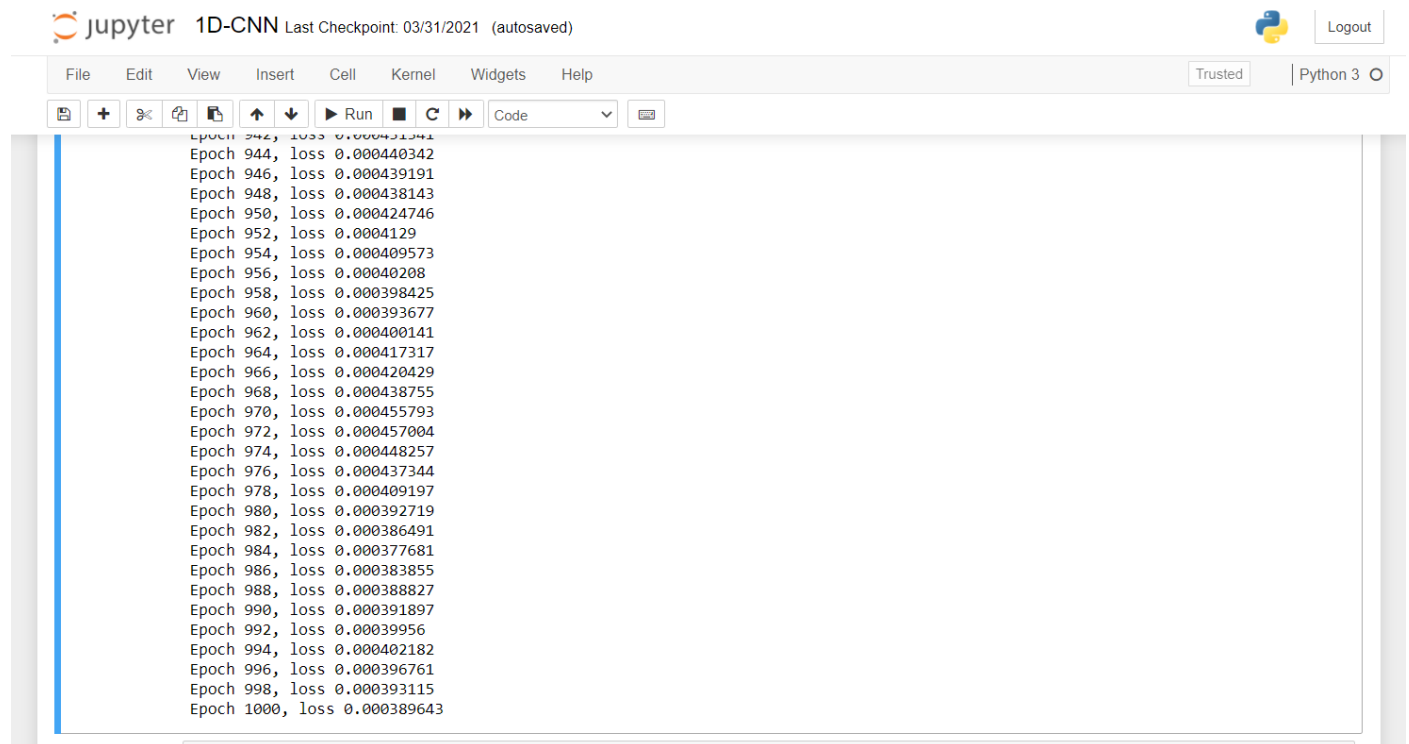
#Once all the epochs are completed, training is stopped
if count >= num_epochs:
    flag = False

count+=1

Epoch 0, loss 0.0127375
Epoch 2, loss 0.00856482
Epoch 4, loss 0.00694965
Epoch 6, loss 0.00555794
Epoch 8, loss 0.00429112
Epoch 10, loss 0.00355224
Epoch 12, loss 0.00311731
Epoch 14, loss 0.00281695
Epoch 16, loss 0.0025776
Epoch 18, loss 0.00236361

C:\Users\hp\anaconda3\lib\site-packages\tensorflow\python\keras\legacy_tf_layers\convolutional.py:202: UserWarning: `tf.layers.conv1d` is deprecated and will be removed in a future version. Please use `tf.keras.layers.Conv1D` instead.
warnings.warn("`tf.layers.conv1d` is deprecated and ")
C:\Users\hp\anaconda3\lib\site-packages\tensorflow\python\keras\engine\base_layer_v1.py:1719: UserWarning: `layer.apply` is deprecated and will be removed in a future version. Please use `layer.__call__` method instead.
warnings.warn("`layer.apply` is deprecated and ")
C:\Users\hp\anaconda3\lib\site-packages\tensorflow\python\keras\legacy_tf_layers\pooling.py:160: UserWarning: `tf.layers.max_pooling1d` is deprecated and will be removed in a future version. Please use `tf.keras.layers.MaxPooling1D` instead.
warnings.warn("`tf.layers.max_pooling1d` is deprecated and ")
C:\Users\hp\anaconda3\lib\site-packages\tensorflow\python\keras\legacy_tf_layers\core.py:329: UserWarning: `tf.layers.flatten` is deprecated and will be removed in a future version. Please use `tf.keras.layers.Flatten` instead.
warnings.warn("`tf.layers.flatten` is deprecated and ")
C:\Users\hp\anaconda3\lib\site-packages\tensorflow\python\keras\legacy_tf_layers\core.py:171: UserWarning: `tf.layers.dense` is deprecated and will be removed in a future version. Please use `tf.keras.layers.Dense` instead.
warnings.warn("`tf.layers.dense` is deprecated and ")
```

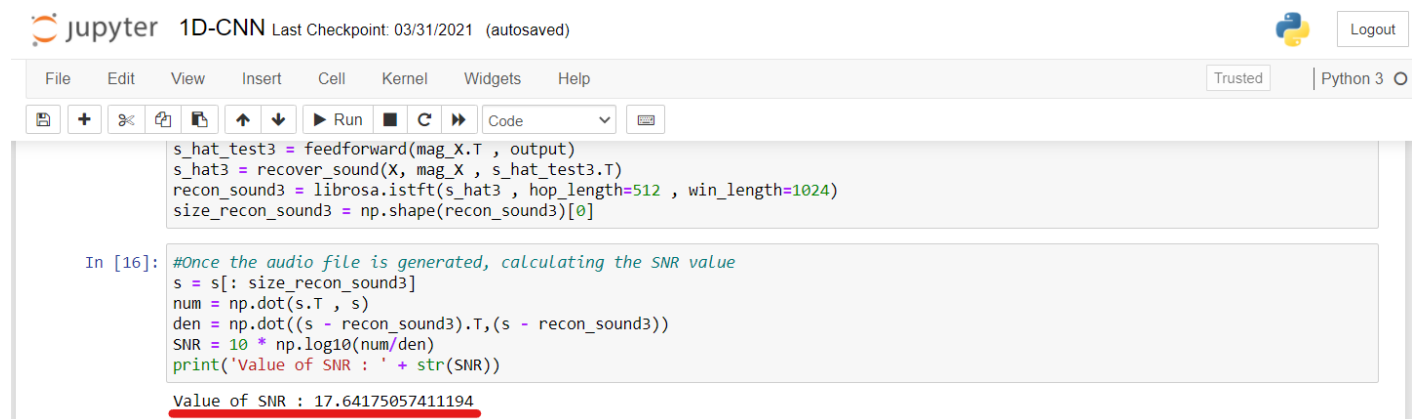
Training completes with 1000 epochs



The image shows a Jupyter Notebook interface with the title "1D-CNN Last Checkpoint: 03/31/2021 (autosaved)". The interface includes a menu bar (File, Edit, View, Insert, Cell, Kernel, Widgets, Help) and a toolbar with icons for saving, adding cells, undo, redo, and running code. The main area displays a list of training epochs and their corresponding loss values. The loss values are consistently low, indicating successful training.

```
Epoch 942, loss 0.000421341
Epoch 944, loss 0.000440342
Epoch 946, loss 0.000439191
Epoch 948, loss 0.000438143
Epoch 950, loss 0.000424746
Epoch 952, loss 0.0004129
Epoch 954, loss 0.000409573
Epoch 956, loss 0.00040208
Epoch 958, loss 0.000398425
Epoch 960, loss 0.000393677
Epoch 962, loss 0.000400141
Epoch 964, loss 0.000417317
Epoch 966, loss 0.000420429
Epoch 968, loss 0.000438755
Epoch 970, loss 0.000455793
Epoch 972, loss 0.000457004
Epoch 974, loss 0.000448257
Epoch 976, loss 0.000437344
Epoch 978, loss 0.000409197
Epoch 980, loss 0.000392719
Epoch 982, loss 0.000386491
Epoch 984, loss 0.000377681
Epoch 986, loss 0.000383855
Epoch 988, loss 0.000388827
Epoch 990, loss 0.000391897
Epoch 992, loss 0.00039956
Epoch 994, loss 0.000402182
Epoch 996, loss 0.000396761
Epoch 998, loss 0.000393115
Epoch 1000, loss 0.000389643
```

SNR calculated after loading, training, processing



The image shows a Jupyter Notebook interface with the title "1D-CNN Last Checkpoint: 03/31/2021 (autosaved)". The interface includes a menu bar (File, Edit, View, Insert, Cell, Kernel, Widgets, Help) and a toolbar with icons for saving, adding cells, undo, redo, and running code. The main area displays a code cell with the following code:

```
s_hat_test3 = feedforward(mag_X.T , output)
s_hat3 = recover_sound(X, mag_X , s_hat_test3.T)
recon_sound3 = librosa.istft(s_hat3 , hop_length=512 , win_length=1024)
size_recon_sound3 = np.shape(recon_sound3)[0]

In [16]: #Once the audio file is generated, calculating the SNR value
s = s[: size_recon_sound3]
num = np.dot(s.T , s)
den = np.dot((s - recon_sound3).T,(s - recon_sound3))
SNR = 10 * np.log10(num/den)
print('Value of SNR : ' + str(SNR))

Value of SNR : 17.64175057411194
```

SNR: 17.641750

Noisy Input Files

Noisy audio input file 1

```
In [28]: # Noisy audio file 1 input
import IPython.display as ipd
print("test_x_01 dirty\n")
ipd.Audio('test_x_01.wav')
```

test_x_01 dirty

Out[28]:




Noisy audio input file 2

```
In [32]: # Noisy audio file 2 input
import IPython.display as ipd
print("test_x_02 dirty\n")
ipd.Audio('test_x_02.wav')
```

test_x_02 dirty

Out[32]:




Denoised audio output files by 1D CNN

Denoised audio output file 1 by 1D CNN

```
In [29]: # 1D CNN denoised audio file 1 output
import IPython.display as ipd
print("1D CNN output_test_x_01\n")
ipd.Audio('test_s_01_recons_q1.wav')
```

1D CNN output_test_x_01

Out[29]:



Denoised audio output file 2 by 1D CNN

```
In [33]: # 1D CNN denoised audio file 2 output
import IPython.display as ipd
print("1D CNN output_test_x_02\n")
ipd.Audio('test_s_02_recons_q1.wav')
```

1D CNN output_test_x_02

Out[33]:




Denoised audio output files by 1D CNN

Denoised audio file 1 by 2D CNN

```
In [91]: # 2D CNN denoised audio file 1 output
import IPython.display as ipd
print("2D CNN output_test_x_02\n")
ipd.Audio('test_s_01_recons_q2.wav')
```

2D CNN output_test_x_02

Out[91]:




Denoised audio file 2 by 2D CNN

```
In [92]: # 2D CNN denoised audio file 2 output
import IPython.display as ipd
print("2D CNN output_test_x_02\n")
ipd.Audio('test_s_02_recons_q2.wav')
```

2D CNN output_test_x_02


Out[92]:



Waveforms of audio files (1D CNN)


Waveform of input noisy audio file 1

```
In [23]: # Test noisy audio file 1
rate, data = wav.read('test_x_01.wav')
%matplotlib inline
plt.plot(data)
plt.show()
```



Waveform of output denoised audio file 1 by 1D CNN


```
In [25]: # 1D CNN denoised audio file 1
rate, data = wav.read('test_s_01_recons_q1.wav')
%matplotlib inline
plt.plot(data)
plt.show()
```



Waveforms of audio files (1D CNN)


Waveform of input noisy audio file 2

```
In [34]: # Test noisy audio file 2
rate, data = wav.read('test_x_02.wav')
%matplotlib inline
plt.plot(data)
plt.show()
```



Waveform of output denoised audio file 2 by 1D CNN


```
In [35]: # 1D CNN denoised audio file 2
rate, data = wav.read('test_s_02_recons_q1.wav')
%matplotlib inline
plt.plot(data)
plt.show()
```



Waveforms of audio files (2D CNN)


Waveform of input noisy audio file 1

```
In [23]: # Test noisy audio file 1
rate, data = wav.read('test_x_01.wav')
%matplotlib inline
plt.plot(data)
plt.show()
```



2D CNN Denoised audio file 1


```
In [96]: # 2D CNN denoised audio file 1
rate, data = wav.read('test_s_01_recons_q2.wav')
%matplotlib inline
plt.plot(data)
plt.show()
```



Waveforms of audio files (2D CNN)


Waveform of input noisy audio file 2

```
In [34]: # Test noisy audio file 2
rate, data = wav.read('test_x_02.wav')
%matplotlib inline
plt.plot(data)
plt.show()
```



Waveform of denoised audio file 2 by 2D CNN

```
In [97]: # 2D CNN denoised audio file 2
rate, data = wav.read('test_s_02_recons_q2.wav')
%matplotlib inline
plt.plot(data)
plt.show()
```



2D CNN training

Training started

jupyter

2D-CNN

Last Checkpoint: 26 minutes ago (unsaved changes)

Logout

FileEditViewInsertCellKernelWidgetsHelp

TrustedPython 3

+

%

↑

↓

▶ Run

■

↺

↻

▶▶

Code

▼

```
batch_x = batch_x.reshape((np.shape(batch_x)[0] * np.shape(batch_x)[1] , np.shape(batch_x)[2]))
feed_dict = {input: batch_x, labels: batch_y}
train_step.run(feed_dict=feed_dict)

if count%1 == 0:
    loss_calc = loss.eval(feed_dict=feed_dict)
    print("Epoch %d, loss %g"%(count+1, loss_calc))
```

100% 100/100 [21:24<00:00, 12.84s/it]

100% 39/39 [00:13<00:00, 2.79it/s]

Epoch 1, loss 0.00550564

100% 39/39 [06:30<00:00, 10.01s/it]

Epoch 2, loss 0.00396692

100% 39/39 [00:27<00:00, 1.42it/s]

Epoch 3, loss 0.0029008

100% 39/39 [00:13<00:00, 2.85it/s]

Training ended

FileEditViewInsertCellKernelWidgetsHelp

TrustedPython 3

+

%

↑

↓

▶ Run

■

↺

↻

▶▶

Code

▼

Epoch 497, loss 0.000345829

100% 39/39 [00:26<00:00, 1.46it/s]

Epoch 498, loss 0.000338241



100% 39/39 [00:13<00:00, 2.90it/s]

Epoch 499, loss 0.000332301











100% 39/39 [00:13<00:00, 2.93it/s]

Epoch 500, loss 0.000328994

2D CNN SNR: 14.63251

 **jupyter** 2D-CNN Last Checkpoint: 39 minutes ago (unsaved changes)  Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3

         Code 

```
sf.write('test_s_01_recons_q1.wav', recon_sound, sr)

# data, samplerate = sf.read('existing_file.wav')
recon_sound2 = librosa.istft(s_hat2, hop_length=512, win_length=1024)
sf.write('test_s_02_recons_q1.wav', recon_sound2, sr)

In [86]: #For testing purpose, feeding the model with train_dirty_male file
#From the output generated, reconstructing the audio file
s_hat_test3 = feedforward(transformed_x1, output)
recovered_x1 = recover_data(s_hat_test3, (window_size - 1, np.shape(s_hat_test3)[1]), 1e-15)
s_hat3 = recover_sound(X, mag_X, recovered_x1.T)
recon_sound3 = librosa.istft(s_hat3, hop_length=512, win_length=1024)
size_recon_sound3 = np.shape(recon_sound3)[0]

In [85]: #Once the audio file is generated, calculating the SNR value
#For testing purpose, feeding the model with train_dirty_male file
#From the output generated, reconstructing the audio file
s_hat_test3 = feedforward(transformed_x1, output)
recovered_x1 = recover_data(s_hat_test3, (window_size - 1, np.shape(s_hat_test3)[1]), 1e-15)
s_hat3 = recover_sound(X, mag_X, recovered_x1.T)
recon_sound3 = librosa.istft(s_hat3, hop_length=512, win_length=1024)
size_recon_sound3 = np.shape(recon_sound3)[0]
s = s[:size_recon_sound3]
num = np.dot(s.T, s)
den = np.dot((s - recon_sound3).T, (s - recon_sound3))
SNR = 10 * np.log10(num/den)
print('Value of SNR : ' + str(SNR))


Value of SNR : 14.632517099380493
```

Waveform of 2D CNN Denoised audio file 1

```
In [91]: # 2D CNN denoised audio file 1 output
import IPython.display as ipd
print("2D CNN output_test_x_02\n")
ipd.Audio('test_s_01_recons_q2.wav')
```

2D CNN output_test_x_02

Out[91]:




Waveform of 2D CNN Denoised audio file 2

```
In [92]: # 2D CNN denoised audio file 2 output
import IPython.display as ipd
print("2D CNN output_test_x_02\n")
ipd.Audio('test_s_02_recons_q2.wav')
```

2D CNN output_test_x_02

Out[92]:



5) Block Diagram

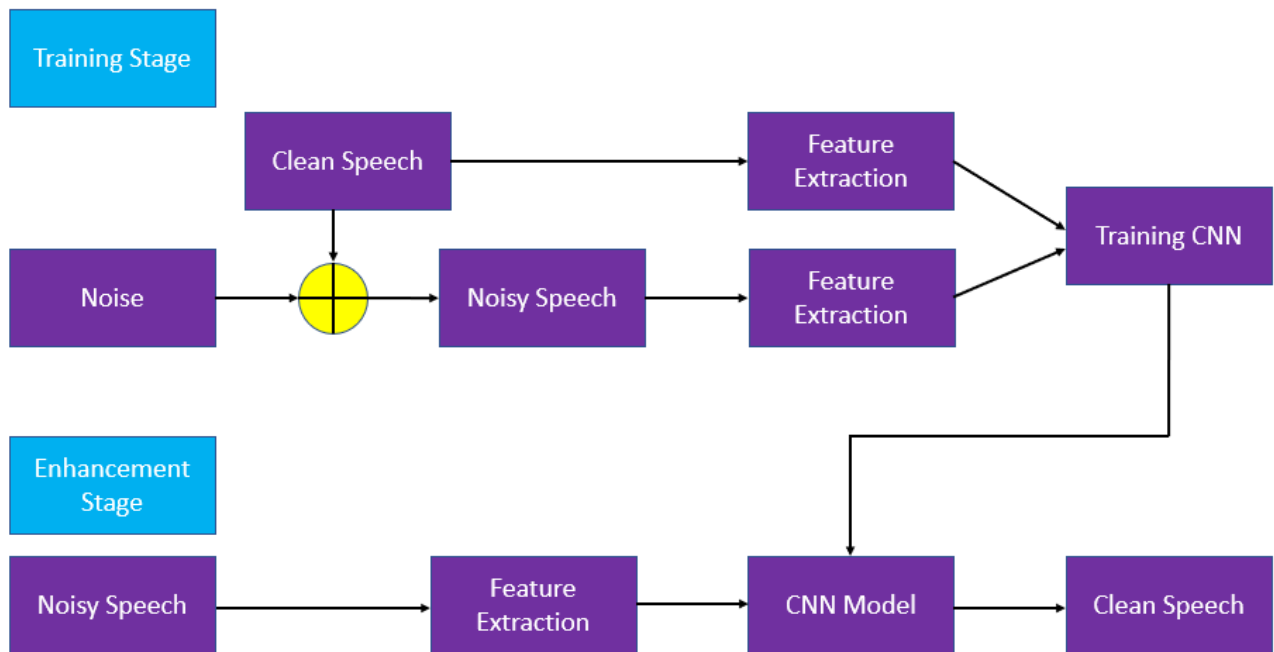


Figure 1: Block Diagram

6) Algorithm Used

- 1) Loading training and testing input audio files using load() function of librosa by which audio files are converted into NumPy array.
- 2) Computing short time fourier transform on all those input audio files array, using stft() function of librosa which transforms this array (consisting amplitude in time domain) to frequency domain. This helps in removing high frequency noise.
- 3) Calculate magnitude of all input files using NumPy.
- 4) Defining CNN model specifications like learning rate, epochs, no. of layers, filters, kernels. After every convolutional layer we are adding pooling layer. Non linearity ReLu has been applied to feature maps output by convolutional layer.
- 5) We defined mean square loss function with Adam optimizer to minimize the loss between expected and model output to train deep learning model
- 6) We use neural networks to find a transformation between the already pre-processed noisy clip and the clean clip.
- 7) Then we performed maxpooling and added a fully connected layer at the end to reduce the dimensionality and get the desired dimensions.
- 8) After training the model, we tested the performance of the model by calculating the Signal-to-Noise Ratio (SNR) value and performed ISTFT to check how the audio sounded.

1D CNN design implemented:

- Two convolution layers with filters 16 and 32 respectively.
- Also kernel sizes of 16, 8 respectively.
- Same padding is used.
- ReLU activation function is used in all the convolution layers.
- Max pooling layers are implemented one each after the convolution layer.
- Flattening is implemented for the last max pooling layer.
- A dense layer of 513 units with a ReLU activation.
- Adam optimizer, mean squared error loss function are used.
- 1000 epochs are used for training.

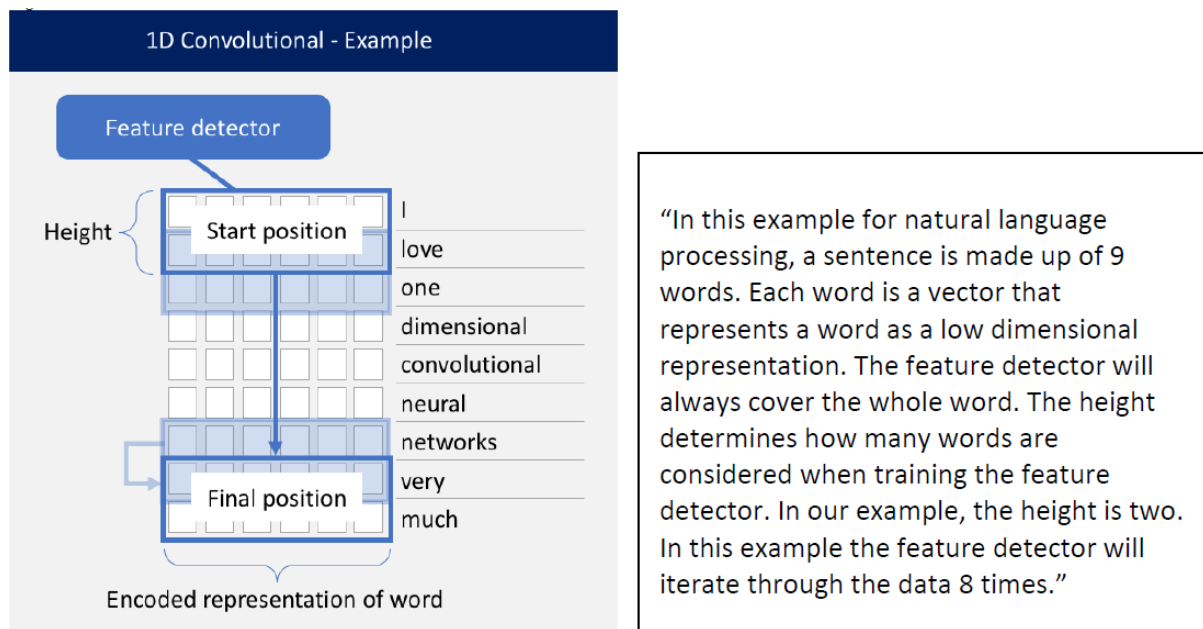


Figure 2: 1D CNN by Blog/Good Audience

2D CNN design implemented:

- Input layer with Tensor Shape as $(-1, 20, 513, 1)$.
- Two convolution layers used with filters of 16, 32 are used respectively.
- Also kernel size of $(4, 4)$ for all of the above layers.
- ReLU activation is used in all the convolution layers.
- Max pooling layers are used after each convolution layer, with `pool_size` of $(2, 2)$.
- Final max pooling layer is flattened.
- A dense layer is used with 513 hidden units with ReLU activation function.
- Adam optimizer (0.0002 learning rate) and Mean squared error Loss function are used.
- Batch size of 64 and 500 epochs for training.

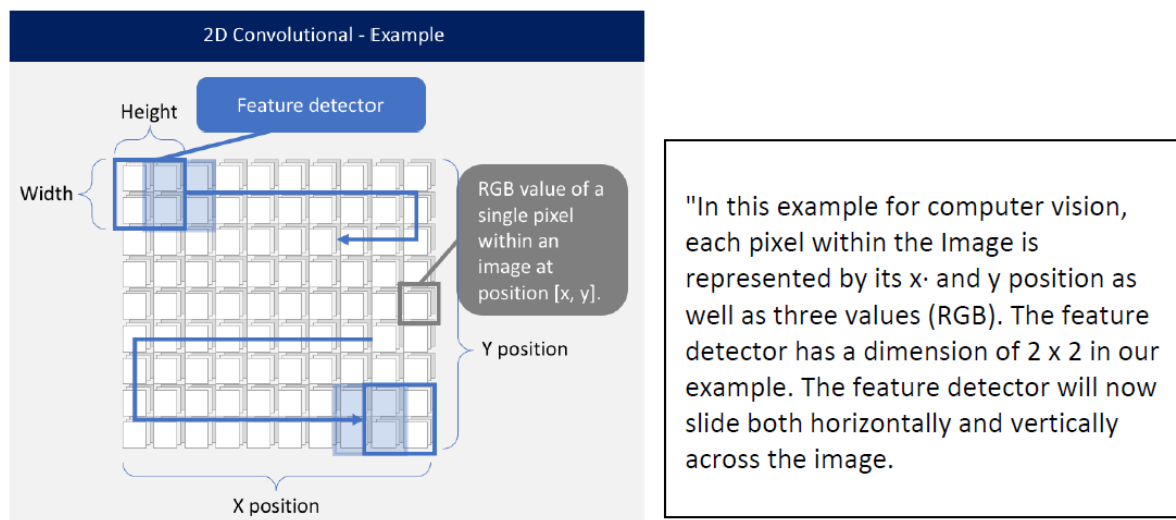


Figure 3: 2D CNN by Blog/Good Audience

7) Observations and Investigations

Comparison

Here, we compare the performances of speech denoising first computing stft then applying deep learning model by 1D CNN & 2D CNN.

Method	Accuracy (SNR)
1D CNN	17.641 DB
2D CNN	14.632 DB

Table 1: Performance of stft + deep learning-based method

8) Conclusion

Hence it is evident from the above results that using (1D CNN + stft) gives better result compared to (2D CNN + stft) for speech denoising by removal of noise from noisy speech and speech enhancement.

By having SNR of 17.64 DB by 1D CNN compared to 14.632 DB by 2D CNN.

- 9) Code Link: <https://github.com/mohdshahbaz123/Speech-Denoising-using-Deep-Learning>